

**YOUR TITLE HERE**

**A THESIS**

*Submitted by*

**YOUR NAME ROLL NO**

*In partial fulfilment for the award of the degree of*

**BACHELOR OF TECHNOLOGY  
IN  
COMPUTER SCIENCE AND ENGINEERING**

**Under the guidance of  
GUIDE NAME**



**DEPARTMENT OF COMPUTER ENGINEERING  
NATIONAL INSTITUTE OF TECHNOLOGY CALICUT  
NIT CAMPUS PO, CALICUT  
KERALA, INDIA 673601**

**April 23, 2015**

## ACKNOWLEDGEMENTS

Your acknowledgements

NAME

## DECLARATION

*“I hereby declare that this submission is my own work and that, to the best of my knowledge and belief, it contains no material previously published or written by another person nor material which has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text”.*

**Place:**  
**Date:**

**Signature :**  
**Name :**  
**Reg.No:**

## CERTIFICATE

*This is to certify that the thesis entitled: “YOUR TITLE HERE” submitted by Sri/Smt/Ms **YOUR NAME** **ROLL NO** to National Institute of Technology Calicut towards partial fulfillment of the requirements for the award of Degree of Bachelor of Technology in Computer Science Engineering is a bonafide record of the work carried out by him/her under my/our supervision and guidance.*

*Signed by Thesis Supervisor(s) with name(s) and date*

**Place:**

**Date:**

*Signature of Head of the Department*

*Office Seal*

## Contents

### Chapter

|          |                                     |    |
|----------|-------------------------------------|----|
| <b>1</b> | Problem Definition                  | 1  |
| <b>2</b> | Introduction                        | 2  |
| 2.0.1    | Publish Subscribe Systems . . . . . | 3  |
| 2.1      | Previous Works . . . . .            | 4  |
| 2.1.1    | PASTRY . . . . .                    | 4  |
| 2.1.2    | MEDYM . . . . .                     | 5  |
| 2.1.3    | SIENA . . . . .                     | 6  |
| 2.1.4    | Rebeca . . . . .                    | 7  |
| <b>3</b> | Experimental Setup                  | 9  |
| 3.1      | Simulation Environment . . . . .    | 9  |
|          | <b>Bibliography</b>                 | 11 |
|          | <b>Appendix</b>                     |    |
| <b>A</b> | Appendix                            | 12 |

## **Abstract**

Abstract here. Abstract should not exceed one page. if@

## Tables

### Table

## Figures

### Figure



## Chapter 1

### Problem Definition

Type your problem definition.

## Chapter 2

### Introduction

Today, the architecture of distributed computer systems is dominated by client/ server platforms relying on synchronous request/reply. This architecture is not well suited to implement information-driven applications like news delivery, stock quoting, air traffic control, and dissemination of auction bids due to the inherent mismatch between the demands of these applications and the characteristics of those platforms. In contrast to that, publish/subscribe directly reflects the intrinsic behavior of information-driven applications because communication here is indirect and initiated by producers of information: Producers publish notifications and these are delivered to subscribed consumers by the help of a notification service that decouples the producers and the consumers. Therefore, publish/subscribe should be the first choice for implementing such applications. [?]

In client/server systems two roles exist: A component acts as a client if it requests data or functionality from another component; it acts as a server if it responds to a client's request. Moreover, a client is blocked after it has issued a request, until the corresponding reply arrives. One of the main deficiencies is the tight coupling among the involved components, i.e., the clients and the servers: A client needs to explicitly address the server that shall process the request, the server must be ready and able to process the request, and the client is blocked, until it receives the reply. Because of these inherent disadvantages a large range

of applications cannot be realized efficiently by using request/reply. These problems are approached by a new communication paradigm called publish/subscribe that recently gained increased publicity in the distributed systems research area. Publish/subscribe is an asynchronous communication paradigm that is also the basis for extensions and supplementary services that have been added to standard middleware recently.

### **2.0.1 Publish Subscribe Systems**

A publish/subscribe system consists of a set of clients that asynchronously exchange notifications, decoupled by a notification service that is interposed between them. Clients can be characterized as producers or consumers. Producers publish notifications such as current stock quotes, and consumers subscribe to notifications by issuing subscriptions, which are essentially stateless message filters. Consumers can have multiple active subscriptions, and after a client has issued a subscription the notification service delivers all future matching notifications that are published by any producer until the client cancels the respective subscription. Publish/subscribe systems have a number of interesting characteristics. Firstly, producers do not need to address consumers and vice versa. Secondly, communication is asynchronous, thereby removing the disadvantages and inflexibility of synchronous communication described above. Thirdly, producers and consumers do not need to be available at the same time. Finally, publish/subscribe directly reflects the intrinsic behavior of information-driven applications because communication is initiated by producers of information.

## 2.1 Previous Works

### 2.1.1 PASTRY

Pastry [?] is a scalable, distributed object location and routing substrate for wide-area peer-to-peer applications. Pastry performs application-level routing and object location in a potentially very large overlay network of nodes connected via the Internet.

Each node in the Pastry network has a unique identifier (`nodeId`). When presented with a message and a key, a Pastry node efficiently routes the message to the node with a `nodeId` that is numerically closest to the key, among all currently live Pastry nodes. Each Pastry node keeps track of its immediate neighbors in the `nodeId` space, and notifies applications of new node arrivals, node failures and recoveries. Pastry takes into account network locality; it seeks to minimize the distance messages travel, according to a scalar proximity metric like the number of IP routing hops. Pastry is completely decentralized, scalable, and self-organizing.

|   |  |
|---|--|
| Routing Performance                       | <sup>5</sup><br>logxN where $x=2b$ and N is the total number of nodes.<br>For one it is the order of the diameter of the graph.    |
| No. Of Messages                           | logxN where $x=2b$ and<br>N is the total number of nodes   |
| Routing Table Size                        | leaf set= $2*2b$ and neighbor set= $2*2b$ and<br>routing table= $2b * \log xN$ where $x=2b$<br>and N is the total number of nodes. |
| Overhead on node<br>Addition and deletion | involves passing routing tables among at least logxN nodes<br>(where $x=2b$ and N is the total number of nodes.)                   |
| Fault Tolerance                           | not tolerant   |
| Node Failure                              | self adjusts   |
| Link Failure                              | not specified  |
| Traffic adaptability                      | adaptable to scalar traffic metrics  |
| Requirement for<br>global Knowledge       | not required   |
| Scalability                               | Scalable   |
| Correctness of algorithm                  | correct  |

### 2.1.2 MEDYM

MEDYM: Match Early with DYnamic Multicast[?]. Unlike existing approaches, MEDYM does not build static overlay networks for event delivery. Its event delivery process is as shown in Figure 4. When an event is published, it is first matched against subscriptions from remote servers, to obtain a destination list of servers with matching subscriptions. Then, the event is routed to the destination servers through dynamic multicast: On receiving an event message, based on its destination list, a server dynamically computes the next-hop servers to which to forward the message, as well as the new destination list for each of the next-hop

servers. In this way, a transient dynamic multicast tree is constructed on the fly.

Along this tree, the event is routed to all the servers with matching subscriptions.

|   |   |
|---|---|
| Routing Performance                       | For one subscriber it is the order of the diameter of the graph   |
| No. Of Messages                           | For one subscriber it is the order of the diameter of the graph   |
| Routing Table Size                        | directly proportional to the number of nodes in the network   |
| Overhead on node<br>Addition and deletion | dynamically adjusts to new node arrivals and node deletions   |
| Fault Tolerance                           | tolerance is not part of MEDYM  |
| Node Failure                              | The topological changes do not affect the algorithm   |
| Link Failure                              | The topological changes do not affect the algorithm   |
| Traffic adaptability                      | The routing can be done to optimize any metric  |
| Requirement for<br>global Knowledge       | The complete knowledge of the entire network is required.<br>And the information has to be updated periodically |
| Scalability                               | Not Scalable  |
| Correctness of algorithm                  | correct   |

### 2.1.3 SIENA

SIENA(Scalable Internet Event Notification Architecture)[?]. SIENA's topology is static. The subscriptions are propagated along the the topology to the publisher. and the subscriptions are stored along the path. The publications are delivered in the reverse path.

|   |   |
|---|---|
| Routing Performance                       | For one subscription it is the diameter of the tree. <sup>7</sup>                               |
| No. Of Messages                           | For one subscriber it is the order of the diameter of the graph                                 |
| Routing Table Size                        | proportional to the number of subscriptions in the descendents                                  |
| Overhead on node<br>Addition and deletion | the node addition and deletion has to<br>ensure that the static overlay topology is maintained. |
| Fault Tolerance                           | The system is not adaptable to faults   |
| Node Failure                              | The node failure has is not handled by the SIENA  |
| Link Failure                              | Link failure can jeopardize the system<br>as there are no redundant paths.                      |
| Traffic adaptability                      | The optimization depending upon various metrics is<br>not possible due to static overlay        |
| Requirement for<br>global Knowledge       | not required  |
| Scalability                               | Scalable  |
| Correctness of algorithm                  | correct   |

#### 2.1.4 Rebeca

Rebeca(Rebeca Event Based Electronic Commerce Architecture) also uses a static architecture. The basic routing is similar to SIENA but the subscriptions are stored using covering and merging techniques.

|  |   |
|--|---|
| Routing Performance                    | For one subscription it is the diameter of the tree. <sup>8</sup>   |
| No. Of Messages                        | For one subscriber it is the order of the diameter of the graph   |
| Routing Table Size                     | Subscriptions are stored in a different format.<br>Uses covering, merging techniques to store subscriptions   |
| Overhead on node addition and deletion | should maintain the properties of the static topology<br>overhead depends on the choice of the node.<br>depends on the number of nodes already present in the network |
| Fault Tolerance                        | The faults are tolerated by leased subscriptions.   |
| Node Failure                           | no specific details   |
| Link Failure                           | cannot adapt to link failures   |
| Traffic adaptability                   | cannot adapt to traffic due to static topology  |
| Requirement for global Knowledge       | not required  |
| Scalability                            | Scalable  |
| Correctness of algorithm               | correct   |



## Chapter 3

### Experimental Setup

#### 3.1 Simulation Environment

The Performance evaluation of the algorithms is done on a simulator. The simulator is developed in Java.

The simulator allows the user to create a logical network with nodes and links. The links are of one of the three types( secure ,urgent,normal).The simulation of the network is done by simulation events.( node,link addition and failures,messages).The simulation module contains the classes which processes the simulation events. Each event is processed as a separate thread. the events are processed in the order of the global time stamp which is attached to each event.

The logical network module implements the functionality of the logical node and logical link. This module contains several interfaces which allows different implementations of nodes and links.Routing is also a part of logical network module.The routing Interface provides the flexibility to define different routing algorithms.With each broker node zero or more clients can be associated. EventBroker and Event-Client modules define the respective functionalities.The clients can have two kinds of roles either a publishing role or a subscribing role. The roles module defines these roles. The subscriptions are specified as the filter (Expression on the attributes of the Publishing event).The applnEvents contains the classes that process the filters and checks for matches.Finally the evaluation module measures the performance

in terms of routing table size, node processing load and data traffic and control traffic.

## Bibliography

## Appendix A

### Appendix